Guidelines for the validation of machine learning predictions of species interactions

[Timothée Poisot](#) [1,2]

[1] Université de Montréal   [2] Québec Centre for Biodiversity Sciences

Correspondance to:

Timothée Poisot — `timothee.poisot@umontreal.ca`

1. The prediction of species interactions is gaining momentum as a way to circumvent limitations in data volume. Yet, ecological networks are challenging to predict because they are typically small and sparse. Dealing with extreme class imbalance is a challenge for most binary classifiers, and there are currently no guidelines as to how predictive models can be trained for this specific problem.

2. Using simple mathematical arguments and numerical experiments in which a variety of classifiers (for supervised learning) are trained on simulated networks, we develop a series of guidelines related to the choice of measures to use for model selection, and the degree of unbiasing to apply to the training dataset.

3. Neither classifier accuracy nor the ROC-AUC are informative measures for the performance of interaction prediction. PR-AUC is a fairer assessment of performance. In some cases, even standard measures can lead to selecting a more biased classifier because the effect of connectance is strong. The amount of correction to apply to the training dataset depends on network connectance, on the measure to be optimized, and only weakly on the classifier.

4. These results reveal that training machines to predict networks is a challenging task, and that in virtually all cases, the composition of the training set needs to be experimented on before performing the actual training. We discuss these consequences in the context of the low volume of data.

Ecological networks are a backbone for key ecological and evolutionary processes; yet enumerating all of the interactions they contain is a daunting task, as it scales with S^2, *i.e.* the squared species richness TODO. Recent contributions to the field of ecological network prediction (Pichler et al., 2020; Strydom et al., 2021; **Becker2021OptPre?**) highlight that although interactions can be predicted by adding ecologically relevant information (in the form of, *e.g.* traits), we do not have robust guidelines as to how the predictive ability of these models should be evaluated, nor about how the models should be trained. Here, by relying on simple derivations and a series of simulations, we formulate a number of such guidelines, specifically for the case of binary classifiers derived from thresholded values. Specifically, we conduct an investigation of the models in terms of their skill (ability to make the right prediction), bias (trends towards systematically over-predicting one class), class imbalance (the relative number of cases representing interactions), and show how these effects interact. We conclude on the fact that models with the best interaction-scale predictive score do not necessarily result in the most accurate representation of the network.

The prediction of ecological interactions shares conceptual and methodological issues with two fields in biology: species distribution models (SDMs), and genomics. SDMs suffers from issues affecting interactions prediction, namely low prevalence (due to sparsity of observations/interactions) and data aggregation (due to bias in sampling some locations/species). In previous work, Allouche et al. (2006) suggested that κ was a better test of model performance than the True Skill Statistic (TSS; which we refer to as Youden's informedness); these conclusions were later criticized by Somodi et al. (2017), who emphasized that informedness' is affected both by prevalence and bias. Although this work offers recommendations about the comparison of models, it doesn't establishes baselines or good practices for training on imbalanced ecological data, or ways to remedy the imbalance. Steen et al. (2021) show that, when applying spatial thinning (a process that has no analogue in networks), the best approach to train ML-based SDMs varies according to the balancing of the dataset, and the evaluation measures used. This suggests that there is no single "recipe" that is guaranteed to give the best model. By contrast to networks, SDMs have the advantage of being able to both thin datasets to remove some of the sampling bias (*e.g.* Inman et al., 2021), but also to create pseudo-absences to inflate the number of supposed negatives in the dataset (*e.g.* Iturbide et al., 2015).

An immense body of research on machine learning application to life sciences is focused on genomics (which has very specific challenges, see a recent discussion by Whalen et al., 2021); this sub-field has

generated recommendations that do not necessarily match the current best-practices for SDMs, and therefore hint at the importance of domain-specific guidelines. Chicco & Jurman (2020) suggest using Matthews correlation coefficient (MCC) over F_1, as a protection against over-inflation of predicted results; Delgado & Tibau (2019) advocate against the use of Cohen's κ, again in favor of MCC, as the relative nature of κ means that a worse classifier can be picked over a better one; similarly, Boughorbel et al. (2017) recommend MCC over other measures of performance for imbalanced data, as it has more desirable statistical properties. More recently, Chicco et al. (2021) temper the apparent supremacy of the MCC, by suggesting it should be replaced by Youden's informedness (also known as J, bookmaker's accuracy, and the True-Skill Statistic) when the imbalance in the dataset may not be representative of the actual imbalance.

Species interaction networks are often under-sampled (Jordano, 2016b, 2016a), and this under-sampling is structured taxonomically (Beauchesne et al., 2016), structurally (de Aguiar et al., 2019) and spatially (Poisot, Bergeron, et al., 2021; Wood et al., 2015). As a consequence, networks suffer from data deficiencies both within and between datasets. This implies that the comparison of classifiers across space, when undersampling varies locally (see *e.g.* McLeod et al., 2021) is non-trivial. Furthermore, the baseline value of classifiers performance measures under various conditions of skill, bias, and prevalence, has to be identified to allow researchers to evaluate whether their interaction prediction model is indeed learning. Taken together, these considerations highlight three specific issues for ecological networks. First, what values of performance measures are indicative of a classifier with no skill? This is particularly important as it can evaluate whether low prevalence can lull us into a false sense of predictive accuracy. Second, independently of the question of model evaluation, is low prevalence an issue for *training* or *testing*, and can we remedy it? Finally, because the low amount of data on interaction makes a lot of imbalance correction methods (see *e.g.* Branco et al., 2015) hard to apply, which indicators can be optimized by sacrificing least amount of positive interaction data?

It may sound counter-intuitive to care so deeply about how good a classifier with no-skill is, as by definition, is has no skill. The necessity of this exercise has its roots in the paradox of accuracy: when the desired class ("two species interact") is rare, a model that gets less ecologically performant by only predicting the opposite class ("these two species do not interact") sees its accuracy increase; because most of the guesses have "these two species do not interact" as a correct answer, a model that never predicts interactions would be right an overwhelming majority of the time; it would also be utterly useless. Herein

lies the core challenge of predicting species interactions: the extreme imbalance between classes makes the training of predictive models difficult, and their validation even more so as we do not reliably know which negatives are true. The connectance (the proportion of realized interactions, usually the number of interactions divided by the number of species pairs) of empirical networks is usually well under 20%, with larger networks having a lower connectance (MacDonald et al., 2020), and therefore being increasingly difficult to predict.

A primer on binary classifier evaluation

Binary classifiers, which it to say, machine learning algorithms whose answer is a categorical value, are usually assessed by measuring properties of their confusion matrix, *i.e.* the contingency table reporting true/false positive/negative hits. A confusion matrix is laid out as

$$\begin{pmatrix} \text{tp} & \text{fp} \\ \text{fn} & \text{tn} \end{pmatrix}.$$

In this matrix, tp is the number of times the model predicts an interaction that exists in the network (true positive), fp is the number of times the model predicts an interaction that does not exist in the network (false positive), fn is the number of times the model fails to predict an interaction that actually exists in the network (false negatives), and tn is the number of times the model correctly predicts that an interaction does not exist (true negatives). From these values, we can derive a number of measures of model performance (see Strydom et al., 2021 for a review of their interpretation in the context of networks). At a coarse scale, a classifier is *accurate* when the trace of the matrix divided by the sum of the matrix is close to 1, with other measures informing us on how the predictions fail.

There is an immense diversity of measures to evaluate the performance of classification tasks (Ferri et al., 2009). Here we will focus on five of them with high relevance for imbalanced learning (He & Ma, 2013). The choice of metrics with relevance to class-imbalanced problems is fundamental, because as Japkowicz (2013) unambiguously concluded, "relatively robust procedures used for unskewed data can break down miserably when the data is skewed." Following Japkowicz (2013), we focus on two ranking metrics (the areas under the Receiver Operating Characteristic and Precision Recall curves), and three threshold

metrics (κ, informedness, and MCC; we will briefly discuss F_1 but show early on that it has undesirable properties).

The κ measure (Landis & Koch, 1977) establishes the extent to which two observers (the network and the prediction) agree, and is measured as

$$2\frac{tp \times tn - fn \times fp}{(tp + fp) \times (fp + tn) + (tn + fp) \times (tn + fn)}.$$

Informedness (Youden, 1950) (also known as bookmaker informedness or the True Skill Statistic) is TPR + TNR − 1, where TPR = $tp/(tp + fn)$ and TNR = $tn/(tn + fp)$. Informedness can be used to find the optimal cutpoint in thresholding analyses (Schisterman et al., 2005); indeed, the maximal informedness corresponds to the point on the ROC curve that is closest to the perfect classifier point. The formula for informedness is

$$\frac{tp}{tp + fn} + \frac{tn}{tn + fp} - 1.$$

The MCC is defined as

$$\frac{tp \times tn - fn \times fp}{\sqrt{(tp + fp) \times (tp + fn) \times (tn + fp) \times (tn + fn)}}.$$

Finally, F_1 is the harmonic mean of precision (the chance that interaction was correctly detected as such) and sensitivity (the ability to correctly classify interactions), and is defined as

$$2\frac{tp}{2 \times tp + fp + fn}.$$

A lot of binary classifiers are built by using a regressor (whose task is to guess the value of the interaction, and can therefore return a value considered to be a pseudo-probability); in this case, the optimal value below which predictions are assumed to be negative (*i.e.* the interaction does not exist) can be determined by picking a threshold maximizing some value on the ROC or the PR curve. The area under these curves (ROC-AUC and PR-AUC henceforth) give ideas on the overall goodness of the classifier, and the ideal threshold is the point on these curves that minimizes the tradeoff represented in these curves. Saito &

Rehmsmeier (2015) established that the ROC-AUC is biased towards over-estimating performance for imbalanced data; on the contrary, the PR-AUC is able to identify classifiers that are less able to detect positive interactions correctly, with the additional advantage of having a baseline value equal to prevalence. Therefore, it is important to assess whether these two measures return different results when applied to ecological network prediction. The ROC curve is defined by the false positive rate on the x axis, and the true positive rate on the y axis, and the PR curve is defined by the true positive rate on the x axis, and the positive predictive value on the y axis. By comparison with the previous paragraph, it is obvious that F_1 and MCC have ties to the PR curve (being close to the expected PR-AUC), and that informedness has ties to the ROC curve (whereby the threshold maximizing informedness is also the point of maximal inflection on the ROC curve). One important difference between ROC and PR is that the later does not prominently account for the size of the true negative compartments: in short, it is more sensitive to the correct positive predictions. In a context of strong imbalance, PR-AUC is therefore a more stringent test of model performance.

Baseline values for the threshold metrics

In this section, we will assume a network of connectance ρ, *i.e.* having ρS^2 interactions (where S is the species richness), and $(1 - \rho)S^2$ non-interactions. Therefore, the vector describing the *true* state of the network (assumed to be an unweighted, directed network) is a column vector $\mathbf{o}^T = [\rho, \quad (1 - \rho)]$ (we can safely drop the S^2 terms, as we will work on the confusion matrix, which ends up expressing *relative* values). We will apply skill and bias to this matrix, and measure how a selection of performance metrics respond to changes in these values, in order to assess their suitability for model evaluation.

Confusion matrix with skill and bias

In order to write the values of the confusion matrix for a hypothetical classifier, we need to define two characteristics: its skill, and its bias. Skill, here, refers to the propensity of the classifier to get the correct answer (*i.e.* to assign interactions where they are, and to not assign them where they are not). A no-skill classifier guesses at random, *i.e.* it will guess interactions with a probability ρ. The predictions of a no-skill classifier can be expressed as a row vector $\mathbf{p} = [\rho(1 - \rho)]$. The confusion matrix $\mathbf{M}$ for a no-skill classifier is given by the element-wise (Hadamard, outer) product of these vectors $\mathbf{o} \odot \mathbf{p}$, *i.e.*

$$\mathbf{M} = \begin{pmatrix} \rho^2 & \rho(1-\rho) \\ (1-\rho)\rho & (1-\rho)^2 \end{pmatrix}.$$

In order to regulate the skill of this classifier, we can define a skill matrix $\mathbf{S}$ with diagonal elements equal to s, and off-diagonal elements equal to $(1 - s)$, and re-express the skill-adjusted confusion matrix as $\mathbf{M} \odot \mathbf{S}$, *i.e.*

$$\begin{pmatrix} \rho^2 & \rho(1-\rho) \\ (1-\rho)\rho & (1-\rho)^2 \end{pmatrix} \odot \begin{pmatrix} s & (1-s) \\ (1-s) & s \end{pmatrix}.$$

When $s = 0$, $\mathrm{Tr}(\mathbf{M}) = 0$ (the classifier is *always* wrong), when $s = 0.5$, the classifier is no-skill and guesses at random, and when $s = 1$, the classifier is perfect.

The second element we can adjust in this hypothetical classifier is its bias, specifically its tendency to over-predict interactions. Like above, we can do so by defining a bias matrix $\mathbf{B}$, where interactions are over-predicted with probability b, and express the final classifier confusion matrix as $\mathbf{M} \odot \mathbf{S} \odot \mathbf{B}$, *i.e.*

$$\begin{pmatrix} \rho^2 & \rho(1-\rho) \\ (1-\rho)\rho & (1-\rho)^2 \end{pmatrix} \odot \begin{pmatrix} s & (1-s) \\ (1-s) & s \end{pmatrix} \odot \begin{pmatrix} b & b \\ (1-b) & (1-b) \end{pmatrix}.$$

The final expression for the confusion matrix in which we can regulate the skill and the bias is

$$\mathbf{C} = \begin{pmatrix} s \times b \times \rho^2 & (1-s) \times b \times \rho(1-\rho) \\ (1-s) \times (1-b) \times (1-\rho)\rho & s \times (1-b) \times (1-\rho)^2 \end{pmatrix}.$$

In all further simulations, the confusion matrix $\mathbf{C}$ is transformed so that it sums to unity, *i.e.* the entries are the *proportions* of guesses.

What are the baseline values of performance measures?

In this section, we will change the values of b, s, and ρ, and report how the main measures discussed in the introduction (MCC, F_1, κ, and informedness) respond. Before we do so, it is important to explain why

144 we will not focus on accuracy too much. Accuracy is the number of correct predictions ($\text{Tr}(\mathbf{C})$) divided by

145 the sum of the confusion matrix. For a no-skill, no-bias classifier, accuracy is equal to $\rho^2 + (1 - \rho)^2$; for

146 $\rho = 0.05$, this is ≈ 0.90, and for $\rho = 0.01$, this is equal to ≈ 0.98. In other words, the values of accuracy are

147 high enough to be uninformative (for ρ small, $\rho^2 \ll (1 - \rho)^2$). More concerning is the fact that introducing

148 bias changes the response of accuracy in unexpected ways. Assuming a no-skill classifier, the numerator

149 of accuracy becomes $b\rho^2 + (1 - b)(1 - \rho)^2$, which increases when b is low, which specifically means that at

150 equal skill, a classifier that under-predicts interactions will have higher accuracy than an un-biased

151 classifier (because the value of accuracy is dominated by the size of tn, which will increase). These issues

152 are absent from balanced accuracy, but should nevertheless lead us to not report accuracy as the primary

153 measure of network prediction success; moving forward, we will focus on other measures.

154 In order to examine how MCC, F_1, κ, and informedness change w.r.t. the imbalance, skill, and bias, we

155 performed a grid exploration of the values of $\text{logit}(s)$ and $\text{logit}(b)$ linearly from -10 to 10; $\text{logit}(x) = -10$

156 means that x is essentially 0, and $\text{logit}(x) = 10$ means it is essentially 1 – this choice was motivated by the

157 fact that most responses are non-linear with regards to bias and skill. The values or ρ were taken linearly

158 in $]0, 0.5]$, which is within the range of connectance for species interaction networks. Note that at this

159 point, there is no network model to speak of; the confusion matrix we discuss can be obtained for any

160 classification task. Based on the previous discussion, the desirable properties for a measure of classifier

161 success should be: an increase with classifier skill, especially at low bias; a hump-shaped response to bias,

162 especially at high skill, and ideally centered around $\text{logit}(b) = 0$; an increase with prevalence up until

163 equiprevalence is reached.

164   [Figure 1 about here.]

165 In fig. 1, we show that none of the four measures satisfy all the considerations at once: F_1 increases with

166 skill, and increases monotonously with bias; this is because F_1 does not account for true negatives, and the

167 increase in positive detection masks the over-prediction of interactions. Informedness varies with skill,

168 reaching 0 for a no-skill classifier, but is entirely unsensitive to bias. Both MCC and κ have the same

169 behavior, whereby they increase with skill. κ peaks at increasing values of biass for increasing skill, *i.e.* is

170 likely to lead to the selection of a classifier that over-predicts interactions. By contract, MCC peaks at the

171 same value, regardless of skill, but this value is not $\text{logit}(b) = 0$: unless at very high classifier skill, MCC

172 risks leading to a model that over-predicts interactions. In fig. 2, we show that all measures except F_1 give

a value of 0 for a no-skill classifier, and are forced towards their correct maximal value when skill changes
(*i.e.* a more connected networks will have higher values for a skilled classifierd, and lower values for a
classifier making mostly mistakes).

[Figure 2 about here.]

These two analyses point to the following recommendations: MCC is indeed more appropriate than κ, as
although sensitive to bias, it is sensitive in a consistent way. Informedness is appropriate at discriminating
between different skills, but confounded by bias. As both of these measures bring valuable information on
the model behavior, we will retain them for future analyses. F_1 is increasing with bias, and should not be
prioritized to evalue the performance of the model. The discussion of sensitivity to bias should come with
a domain-specific caveat: although it is likely that interactions documented in ecological networks are
correct, a lot of non-interactions are simply unobserved; as predictive models are used for data-inflation
(*i.e.* the prediction of new interactions), it is not necessarily a bad thing in practice to select models that
predict more interactions than the original dataset, because the original dataset misses some interactions.
Furthermore, the weight of positive interactions could be adjusted if some information about the extent of
undersampling exists (*e.g.* Branco et al., 2015). In a recent large-scale imputation of interactions in the
mammal-virus networks, Poisot, Ouellet, et al. (2021) for example estimated that 93% of interactions are
yet to be documented.

Numerical experiments on training strategy

In the following section, we will generate random bipartite networks (this works without loss of generality
on unipartite networks), and train four binary classifiers (as well as an ensemble model using the sum of
ranged outputs from the component models) on 30% of the interaction data. In practice, testing usually
uses 70% of the total data; for ecological networks, where interactions are sparse *and* the number of
species is low, this may not be the best solution, as the testing set becomes constrained not by the
proportion of interactions, but by their *number*. Preliminary experiments using different splits revealed no
qualitative change in the results. Networks are generated by picking a random infectiousness trait v_i for
100 species (from a beta distribution $B(\alpha = 6, \beta = 8)$ distribution), and a resistance trait h_j for 100 species
(from $B(\alpha = 2, \beta = 8)$ distribution). There is an interaction between i and j when

$v_i - \xi/2 \leq h_j \leq v_i + \xi/2$, where ξ is a constant regulating the connectance of the network (there is an almost 1:1 relationship between ξ and connectance), and varies uniformly in $[0.05, 0.35]$. This model gives fully interval networks that are close analogues to the bacteria–phage model of Weitz et al. (2005), with both a modular structure and a non-uniform degree distribution. This dataset is easy for almost any algorithm to learn: when trained with features $[v_i, h_j, \mathrm{abs}(v_i, h_j)]^T$ to predict the interactions between i and j, all four models presented below were able to reach almost perfect predictions all the time (data not presented here) – this is in part because the rule (there is maximum value of the distance between traits for which there is an interaction) is fixed for all interactions. In order to make the problem more difficult to solve, we use $[v_i, h_j]$ as a feature vector (*i.e.* the traits on which the models are trained), and therefore the models will have to uncover that the rule for interaction is $\mathrm{abs}(v_i, h_j) \leq \xi$. The models therefore all have the following form, where $i_{i,j}$ is an interaction from species i to species j:

$$
\begin{bmatrix} i_{1,1} \\ i_{1,2} \\ \vdots \\ i_{m,n-1} \\ i_{m,n} \end{bmatrix} \propto \begin{bmatrix} v_1 & h_1 \\ v_1 & h_2 \\ \vdots & \vdots \\ v_m & h_{n-1} \\ v_m & h_n \end{bmatrix}
$$

The training sample is composed of 30% of the 10^4 possible entries in the network, *i.e.* $n = 3000$. Out of these interactions, we pick a proportion ν (the training set bias) to be positive, so that the training set has νn interactions, and $(1 - \nu)n$ non-interactions. We vary ν uniformly in $]0, 1[$. This allows to evaluate how the measures of binary classification performance respond to artificially rebalanced dataset for a given network connectance. The rest of the dataset ($n = 7000$ pairs of species) is used as a testing set, on which all furher measures are calculated. Note that although the training set is balanced, the testing set is not, and retains (part of) the imbalance of the original data.

The dataset used for numerical experiments is composed of 64000 such (ξ, ν) pairs, on which four machines are trained: a decision tree regressor, a boosted regression tree, a ridge regressor, and a random forest regressor. All models were taken from the `MLJ.jl` package (Blaom et al., 2020; Blaom & Vollmer, 2020) in Julia 1.7 (Bezanson et al., 2017). All machines use the default parameterization; this is an obvious deviation from best practices, as the hyperparameters of any machine require training before its application on a real dataset. As we use 64000 such datasets, this would require 256000 unique instances

of tweaking the hyperparameters, which is not realistic. Therefore, we assume that the default

parameterizations are comparable across networks. All machines return a quantitative prediction, usually

(but not necessarily) in [0, 1], which is proportional (but not necessarily linearly) to the probability of an

interaction between i and j. Nevertheless, the ROC-AUC and PR-AUC (and therefore the thresholds) can

be measured by integrating over the domain of the values return by each machine.

In order to pick the best adjacency matrix for a given trained machine, we performed a thresholding

approach using 500 steps on predictions from the testing set, and picking the threshold that maximized

Youden's informedness. During the thresholding step, we measured the area under the receiver operating

characteristic (ROC-AUC) and precision-recall (PR-AUC) curves, as measures of overall performance over

the range of returned values. We report the ROC-AUC and PR-AUC, as well as a suite of other measures as

introduced in the next section, for the best threshold. The ensemble model was generated by summing the

predictions of all component models on the testing set (ranged in [0, 1]), then put through the same

thresholding process. The complete code to run the simulations is available at `10.17605/OSF.IO/JKEWD`;

the simulations require approx. 5 core days.

After the simulations were completed, we removed all runs (*i.e.* pairs of ξ and ν) for which at least one of

the following conditions was met: the accuracy was 0, the true positive or true negative rates were 0, the

connectance was larger than 0.25. This removes both the obviously failed model runs, and the networks

that are more densely connected compared to the connectance of empirical food webs (and are therefore

less difficult to predict, being less imbalanced; preliminary analyses of data with a connectance larger than

0.3 revealed that all machines reached consistently high performance).

Effect of training set bias on performance

In fig. 3, we present the response of MCC and informedness to (i) five levels of network connectance and

(ii) a gradient of training set bias, for the four component models as well as the ensemble. All models

reached a higher performance on more connected networks, and using more biased training sets (with the

exception of ridge regression, whose informedness decreased in performance with training set bias). In all

cases, informedness was extremely high, which is an expected consequence of the fact that this is the

value we optimized to determine the cutoff. MCC increased with training set bias, although this increase

became less steep with increasing connectance. Interestingly, the ensemble almost always outclassed its

₂₅₂ component models. In a few cases, both MCC and informedness stared decreasing when the training set

₂₅₃ bias got too close to one, which suggests that it is possible to over-correct the imbalance.

₂₅₄ [Figure 3 about here.]

₂₅₅ In fig. 4, we present the same information as fig. 3, this time using ROC-AUC and PR-AUC. ROC-AUC is

₂₅₆ always high, and does not vary with training set bias. On the other hand, PR-AUC shows very strong

₂₅₇ responses, increasing with training set bias. It is notable here that two classifiers that seemed to be

₂₅₈ performing well (Decision Tree and Random Forest) based on their MCC are not able to reach a high

₂₅₉ PR-AUC even at higher connectances. As in fig. 3, the ensemble outperforms its component models.

₂₆₀ [Figure 4 about here.]

₂₆₁ Based on the results presented in fig. 3 and fig. 4, it seems that informedness and ROC-AUC are not

₂₆₂ necessarily able to discriminate between good and bad classifiers (although this result may be an artifact

₂₆₃ for informedness, as it has been optimized when thresholding). On the other hand, MCC and PR-AUC

₂₆₄ show a strong response to training set bias, and may therefore be more useful at model comparison.

₂₆₅ **Required amount of positives to get the best performance**

₂₆₆ The previous results revealed that the measure of classification performance responds both to the bias in

₂₆₇ the training set *and* to the connectance of the network; from a practical point of view, assembling a

₂₆₈ training set requires to withold positive information, which in ecological networks are very scarce (and

₂₆₉ typically more valuable than negatives, on which there is a doubt). For this reason, across all values of

₂₇₀ connectance, we measured the training set bias that maximized a series of performance measures. When

₂₇₁ this value is high, the training set needs to skew more positive in order to get a performant model; when

₂₇₂ this value is about 0.5, the training set needs to be artificially balanced to optimize the model performance.

₂₇₃ These results are presented in fig. 5.

₂₇₄ [Figure 5 about here.]

₂₇₅ The more "optimistic" measures (ROC-AUC and informedness) required a biasing of the dataset from

₂₇₆ about 0.4 to 0.75 to be maximized, with the amount of bias required decreasing only slightly with the

connectance of the original network. MCC and PR-AUC required values of training set bias from 0.75 to almost 1 to be optimized, which is in line with the results of the previous section, *i.e.* they are more stringent tests of model performance. These results suggest that learning from a dataset with very low connectance can be a different task than for more connected networks: it becomes increasingly important to caputre the mechanisms that make an interaction *exist*, and therefore having a slightly more biased training dataset might be beneficial. As connectance increases, the need for biased training sets is less prominent, as learning the rules for which interactions *do not* exist starts gaining importance.

[Figure 6 about here.]

When trained at their optimal training set bias, connectance still had a significant impact on the performance of some machines fig. 6. Notably, Decision Tree, Random Forest, and Ridge Regression had low values of PR-AUC. In all cases, the Boosted Regression Tree was reaching very good predictions (esepcially for connectances larger than 0.1), and the ensemble was almost always scoring perfectly. This suggests that all the models are biased in different ways, and that the averaging in the ensemble is able to correct these biases. We do not expect this last result to have any generality, and provide a discussion of a recent exemple in which the ensemble was performing worse than its components models.

Do better classification accuracy result in more realistic networks?

In this last section, we generate a network using the same model as before, with $S_1, S_2 = 50, 80$ species, a connectance of ≈ 0.16 ($\xi = 0.19$), and a training set bias of 0.7. The prediction made on the complete dataset is presented in fig. 7. Visualizing the results this way highlights the importance of exploratory data analysis: whereas all models return a network with interactions laying mostly on the diagonal (as expected), the Ridge Regression is quite obviously biased. Despite this, we can see that the ensemble is close to the initial dataset.

[Figure 7 about here.]

The trained models were then thresholded (again by optimising informedness), and their predictions transformed back into networks for analysis; specifically, we measured the connectance, nestedness

302 (REF), and modularity (REF). This process was repeated 250 times, and the results are presented in tbl. 1.

303 The random forest model is an interesting instance here: it produces the network that looks the most like

304 the original dataset, despite having a very low PR-AUC, suggesting it hits high recall at the cost of low

305 precision. Although the ensemble was able to reach a very high PR-AUC (and a very high ROC-AUC), this

306 did not necessarily translate into more accurate reconstructions of the structure of the network. This

307 result bears elaborating. Measures of model performance capture how much of the interactions and

308 non-interactions are correctly identified. As long as these predictions are not perfect, some interactions

309 will be predicted at the "wrong" position in the network; these measures cannot describe the structural

310 effect of these mistakes. On the other hand, measures of network structure can have the same value with

311 interactions that fall at drastically different positions; this is in part because a lot of these measures covary

312 with connectance, and in part because as long as these values are not 0 or their respective maximum, there

313 is a large number of network configurations that can have the same value. That ROC-AUC is consistently

314 larger than PR-AUC may be a case of this measure masking models that are not, individually, strong

315 predictors (Jeni et al., 2013).

Table 1: Values of four performance metrics, and three network structure metrics, for 250 independent predictions similar to the ones presented in fig. 7. The values in **bold** indicate the best value for each column (including ties). Because the values have been rounded, values of 1.0 for the ROC-AUC column indicate an average ≥ 0.99.

Model	MCC	Inf.	ROC-AUC	PR-AUC	Conn.	η	Q
Decision tree	0.85	0.92	0.97	0.12	0.21	0.76	0.31
BRT	**0.90**	0.90	0.98	0.86	0.23	0.82	0.27
Random Forest	**0.90**	**0.96**	**1.00**	0.27	**0.20**	**0.72**	**0.32**
Ridge Regression	0.80	0.91	0.95	0.58	0.24	1.0	0.18
Ensemble	0.88	0.94	**1.00**	**0.96**	**0.20**	0.75	0.31
Data					0.18	0.66	0.34

Guidelines for the assesment of network predictive models

317 We establish that due to the low prevalence of interactions, even poor classifiers applied to food web data

318 will reach a high accuracy; this is because the measure is dominated by the accidentally correct

predictions of negatives. On simulated confusion matrices with ranges of imbalance that are credible for ecological networks, MCC had the most desirable behavior, and informedness is a linear measure of classifier skill. By performing simulations with four models and an ensemble, we show that informedness and ROC-AUC are consistently high on network data, and that MCC and PR-AUC are more accurate measures of the effective performance of the classifier. Finally, by measuring the structure of predicted networks, we highlight an interesting paradox: the models with the best performance measures are not the models with the closest reconstructed network structure. We discuss these results in the context of establishing guidelines for the prediction of ecological interactions.

TODO informedness and accuracy should be easy to beat, make sure the model is better than them!

The results presented here highlight an interesting paradox: although the Random Forest was ultimately able to get a correct estimate of network structure tbl. 1, it ultimately remains a poor classifier, as evidenced by its low PR-AUC. This suggests that the goal of predicting *interactions* and predicting *networks* may not be solvable in the same way – of course a perfect classifier of interactions would make a perfect network prediction; but even the best scoring predictor of interactions (the ensemble model) had not necessarily the best prediction of network structure. The tasks of predicting networks structure and of predicting interactions within networks are essentially two different ones. For some applications (.*e.g.* comparison of network structure across gradients), one may care more about a robust estimate of the structure, at the cost at putting some interactions at the wrong place. For other applications (*e.g.* identifying pairs of interacting species), one may conversely care more about getting as many pairs right, even though the mistakes accumulate in the form of a slightly worse estimate of network structure. How these two approaches can be reconciled is undoubtedly a task for further research. Despite this apparent tension at the heart of the predictive exercise, we can use the results presented here to suggest a number of guidelines.

First, because we should have more trust in reported interactions than in reported absences of interactions, we can draw on previous literature to recommend informedness as a measure to decide on a threshold (Chicco et al., 2021); this being said, because informedness is insensitive to bias, the model performance is better evaluated through the use of MCC fig. 3. Because F_1 is monotonously sensitive to classifier bias fig. 1 and network connectance fig. 2, MCC should be prefered as a measure of model evaluation.

Second, because the PR-AUC responds more to network connectance fig. 6 and training set imbalance fig. 4, it should be used as a measure of model performance over the ROC-AUC. This is not to say that

ROC-AUC should be discarded (in fact, a low ROC-AUC is a sign of an issue with the model), but that its interpretation should be guided by the PR-AUC value. Specifically, a high ROC-AUC is not informative, as it can be associated to a low PR-AUC (see *e.g.* Random Forest in tbl. 1) This again echoes recommendations from other fields (Jeni et al., 2013; Saito & Rehmsmeier, 2015).

Thirdly, regardless of network connectance, maximizing informedness required a training set bias of about 0.5, and maximizing the MCC required a training set bias of 0.7 and more. This has an important consequence in ecological networks, for which the pool of positive cases (interactions) to draw from is typically small: the most parsimonious measure (*i.e.* the one requiring to discard the least amount of information to train the model) will give the best validation potential, and is probably the informedness (maximizing informedness is the generally accepted default for imbalanced classification; Schisterman et al., 2005).

Finally, it is noteworthy that the ensemble model was systematically better than the component models; even when the models were individually far form perfect, the ensemble was able to leverage the different biases expressed by the models to make an overall more accurate prediction. We do not expect that ensembles will *always* be better than single models. In a recent multi-model comparison, (**Becker2021OptPre?**) found that the ensemble was *not* the best model. There is no general conclusion to draw from this besides reinforcing the need to be pragmatic about which models should be included in the ensemble, or whether to use an ensemble at all. In a sense, the surprising peformance of the ensemble model should form the basis of the last recommendation: optimal training set bias and its interaction with connectance and binary classifier is, in a sense, an hyperparameter that should be assessed. The distribution of results in fig. 5 and fig. 6 show that there are variations around the trend; furthermore, networks with different structures than the one we simulated here may respond in different ways.

Valorisation des Données (IVADO).

References

Allouche, O., Tsoar, A., & Kadmon, R. (2006). Assessing the accuracy of species distribution models: Prevalence, kappa and the true skill statistic (TSS). *Journal of Applied Ecology*, *43*(6), 1223–1232. https://doi.org/10.1111/j.1365-2664.2006.01214.x

Beauchesne, D., Desjardins-Proulx, Archambault, P., & Gravel, D. (2016). Thinking Outside the Boxpredicting Biotic Interactions in Data-poor Environments. *Vie Et Milieu-Life and enVironment*, *66*(3-4), 333–342.

Bezanson, J., Edelman, A., Karpinski, S., & Shah, V. (2017). Julia: A Fresh Approach to Numerical Computing. *SIAM Review*, *59*(1), 65–98. https://doi.org/10.1137/141000671

Blaom, A. D., Kiraly, F., Lienart, T., Simillides, Y., Arenas, D., & Vollmer, S. J. (2020). MLJ: A Julia package for composable machine learning. *Journal of Open Source Software*, *5*(55), 2704. https://doi.org/10.21105/joss.02704

Blaom, A. D., & Vollmer, S. J. (2020). Flexible model composition in machine learning and its implementation in MLJ. *arXiv:2012.15505 [cs]*. http://arxiv.org/abs/2012.15505

Boughorbel, S., Jarray, F., & El-Anbari, M. (2017). Optimal classifier for imbalanced data using Matthews Correlation Coefficient metric. *PloS One*, *12*(6), e0177678. https://doi.org/10.1371/journal.pone.0177678

Branco, P., Torgo, L., & Ribeiro, R. (2015). A Survey of Predictive Modelling under Imbalanced Distributions. *arXiv:1505.01658 [cs]*. http://arxiv.org/abs/1505.01658

Chicco, D., & Jurman, G. (2020). The advantages of the Matthews correlation coefficient (MCC) over F1 score and accuracy in binary classification evaluation. *BMC Genomics*, *21*(1), 6. https://doi.org/10.1186/s12864-019-6413-7

Chicco, D., Tötsch, N., & Jurman, G. (2021). The Matthews correlation coefficient (MCC) is more reliable than balanced accuracy, bookmaker informedness, and markedness in two-class confusion matrix evaluation. *BioData Mining*, *14*, 13. https://doi.org/10.1186/s13040-021-00244-z

de Aguiar, M. A. M., Newman, E. A., Pires, M. M., Yeakel, J. D., Boettiger, C., Burkle, L. A., Gravel, D., Guimarães, P. R., O'Donnell, J. L., Poisot, T., Fortin, M.-J., & Hembry, D. H. (2019). Revealing biases in the sampling of ecological interaction networks. *PeerJ*, *7*, e7566. https://doi.org/10.7717/peerj.7566

Delgado, R., & Tibau, X.-A. (2019). Why Cohen's Kappa should be avoided as performance measure in classification. *PloS One*, *14*(9), e0222916. https://doi.org/10.1371/journal.pone.0222916

Ferri, C., Hernández-Orallo, J., & Modroiu, R. (2009). An experimental comparison of performance measures for classification. *Pattern Recognition Letters*, *30*(1), 27–38. https://doi.org/10.1016/j.patrec.2008.08.010

He, H., & Ma, Y. (Eds.). (2013). *Imbalanced Learning: Foundations, Algorithms, and Applications* (1st edition). Wiley-IEEE Press.

Inman, R., Franklin, J., Esque, T., & Nussear, K. (2021). Comparing sample bias correction methods for species distribution modeling using virtual species. *Ecosphere*, *12*(3), e03422. https://doi.org/10.1002/ecs2.3422

Iturbide, M., Bedia, J., Herrera, S., del Hierro, O., Pinto, M., & Gutiérrez, J. M. (2015). A framework for species distribution modelling with improved pseudo-absence generation. *Ecological Modelling*, *312*, 166–174. https://doi.org/10.1016/j.ecolmodel.2015.05.018

Japkowicz, N. (2013). Assessment Metrics for Imbalanced Learning. In *Imbalanced Learning* (pp. 187–206). John Wiley & Sons, Ltd. https://doi.org/10.1002/9781118646106.ch8

Jeni, L. A., Cohn, J. F., & De La Torre, F. (2013). Facing Imbalanced DataRecommendations for the Use of Performance Metrics. *2013 Humaine Association Conference on Affective Computing and Intelligent Interaction*, 245–251. https://doi.org/10.1109/ACII.2013.47

Jordano, P. (2016a). Chasing Ecological Interactions. *PLOS Biol*, *14*(9), e1002559. https://doi.org/10.1371/journal.pbio.1002559

Jordano, P. (2016b). Sampling networks of ecological interactions. *Functional Ecology*. https://doi.org/10.1111/1365-2435.12763

Landis, J. R., & Koch, G. G. (1977). The Measurement of Observer Agreement for Categorical Data. *Biometrics*, *33*(1), 159–174. https://doi.org/10.2307/2529310

MacDonald, A. A. M., Banville, F., & Poisot, T. (2020). Revisiting the Links-Species Scaling Relationship in Food Webs. *Patterns*, *1*(0). https://doi.org/10.1016/j.patter.2020.100079

McLeod, A., Leroux, S. J., Gravel, D., Chu, C., Cirtwill, A. R., Fortin, M.-J., Galiana, N., Poisot, T., & Wood, S. A. (2021). Sampling and asymptotic network properties of spatial multi-trophic networks. *Oikos*, *n/a*(n/a). https://doi.org/10.1111/oik.08650

Pichler, M., Boreux, V., Klein, A.-M., Schleuning, M., & Hartig, F. (2020). Machine learning algorithms to infer trait-matching and predict species interactions in ecological networks. *Methods in Ecology and Evolution*, *11*(2), 281–293. https://doi.org/10.1111/2041-210X.13329

Poisot, T., Bergeron, G., Cazelles, K., Dallas, T., Gravel, D., MacDonald, A., Mercier, B., Violet, C., & Vissault, S. (2021). Global knowledge gaps in species interaction networks data. *Journal of Biogeography*, jbi.14127. https://doi.org/10.1111/jbi.14127

Poisot, T., Ouellet, M.-A., Mollentze, N., Farrell, M. J., Becker, D. J., Albery, G. F., Gibb, R. J., Seifert, S. N., & Carlson, C. J. (2021). Imputing the mammalian virome with linear filtering and singular value decomposition. *arXiv:2105.14973 [q-Bio]*. http://arxiv.org/abs/2105.14973

Saito, T., & Rehmsmeier, M. (2015). The Precision-Recall Plot Is More Informative than the ROC Plot When Evaluating Binary Classifiers on Imbalanced Datasets. *PLOS ONE*, *10*(3), e0118432. https://doi.org/10.1371/journal.pone.0118432

Schisterman, E. F., Perkins, N. J., Liu, A., & Bondell, H. (2005). Optimal Cut-point and Its Corresponding Youden Index to Discriminate Individuals Using Pooled Blood Samples. *Epidemiology*, *16*(1), 73–81. https://doi.org/10.1097/01.ede.0000147512.81966.ba

Somodi, I., Lepesi, N., & Botta-Dukát, Z. (2017). Prevalence dependence in model goodness measures with special emphasis on true skill statistics. *Ecology and Evolution*, *7*(3), 863–872. https://doi.org/10.1002/ece3.2654

Steen, V. A., Tingley, M. W., Paton, P. W. C., & Elphick, C. S. (2021). Spatial thinning and class balancing: Key choices lead to variation in the performance of species distribution models with citizen science data. *Methods in Ecology and Evolution*, *12*(2), 216–226. https://doi.org/10.1111/2041-210X.13525

Strydom, T., Catchen, M. D., Banville, F., Caron, D., Dansereau, G., Desjardins-Proulx, P., Forero-Muñoz, N. R., Higino, G., Mercier, B., Gonzalez, A., Gravel, D., Pollock, L., & Poisot, T. (2021). A roadmap

towards predicting species interaction networks (across space and time). *Philosophical Transactions of the Royal Society B: Biological Sciences*, *376*(1837), 20210063. https://doi.org/10.1098/rstb.2021.0063

Weitz, J. S., Hartman, H., & Levin, S. A. (2005). Coevolutionary arms races between bacteria and bacteriophage. *Proceedings of the National Academy of Sciences of the United States of America*, *102*(27), 9535–9540. https://doi.org/10.1073/pnas.0504062102

Whalen, S., Schreiber, J., Noble, W. S., & Pollard, K. S. (2021). Navigating the pitfalls of applying machine learning in genomics. *Nature Reviews Genetics*, 1–13. https://doi.org/10.1038/s41576-021-00434-9

Wood, S. A., Russell, R., Hanson, D., Williams, R. J., & Dunne, J. A. (2015). Effects of spatial scale of sampling on food web structure. *Ecology and Evolution*, *5*(17), 3769–3782. https://doi.org/10.1002/ece3.1640

Youden, W. J. (1950). Index for rating diagnostic tests. *Cancer*, *3*(1), 32–35. https://doi.org/10.1002/1097-0142(1950)3:1%3C32::AID-CNCR2820030106%3E3.0.CO;2-3

Figure 1: Consequences of changing the classifier skills (s) and bias (s) for a connectance $\rho = 0.15$, on accuracy, F_1, postive predictive value, and κ. Accuracy increases with skill, but also increases when the bias tends towards estimating *fewer* interactions. The F_1 score increases with skill but also increases when the bias tends towards estimating *more* interactions; PPV behaves in the same way. Interestingly, κ responds as expected to skill (being negative whenever $s < 0.5$), and peaks for values of $b \approx 0.5$; nevertheless, the value of bias for which κ is maximized in *not* $b = 0.5$, but instead increases with classifier skill. In other words, at equal skill, maximizing κ would lead to select a *more* biased classifier.

Figure 2: As in fig. 1, consequences of changing connectance for different levels of classifier skill, assuming no classifier bias. Informedness, κ, and MCC do increase with connectance, but only when the classifier is not no-skill; by way of contrast, a more connected network will give a higher F_1 value even with a no-skill classifier.
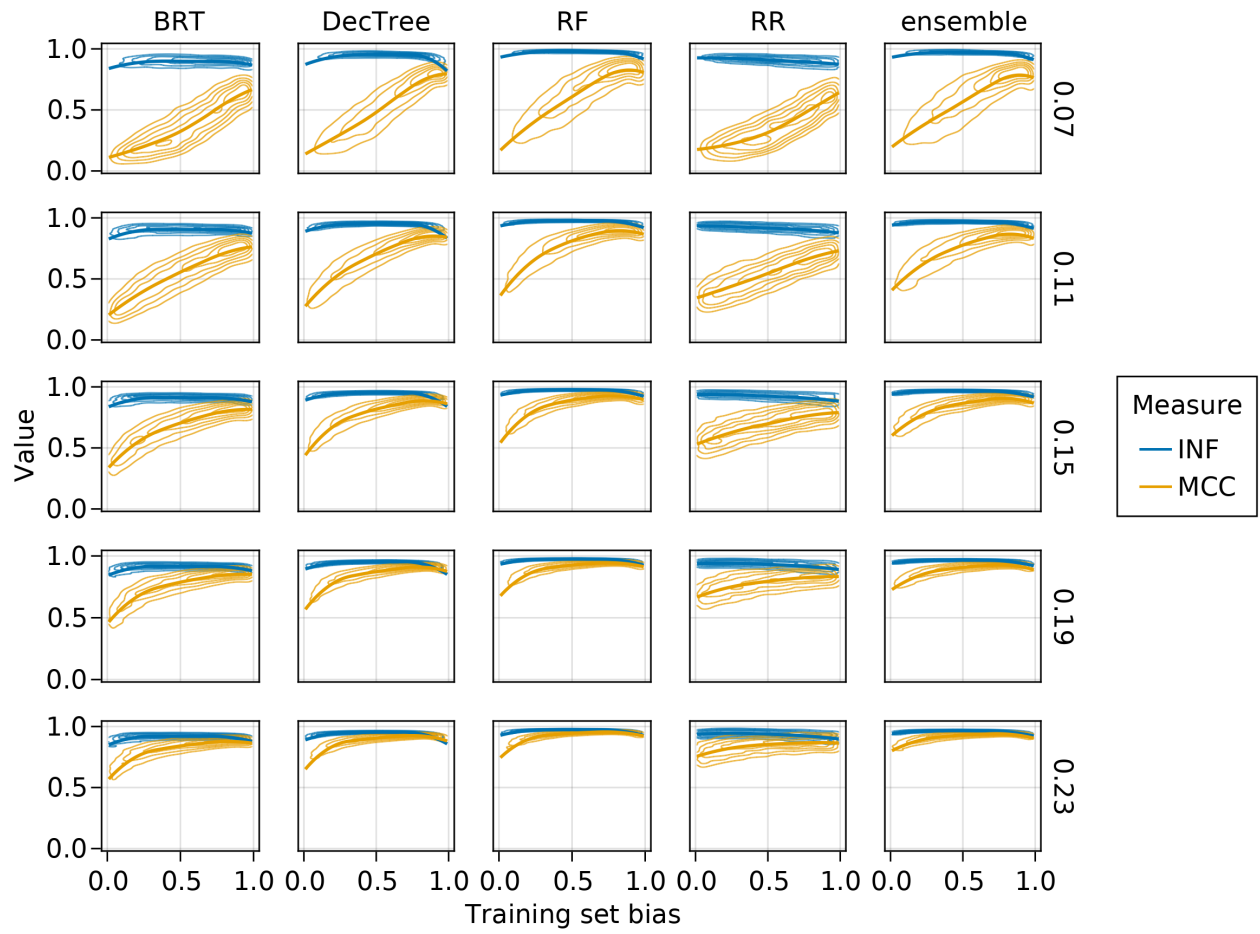
Figure 3: Response of MCC and Informedness to changes in the training set bias for a fixed connectance (rows). Both of these values approach 1 for a good model. Informedness is consistently high, and by contrast, MCC increases with additional training set bias. Across all models, training on a more connected network is easier.
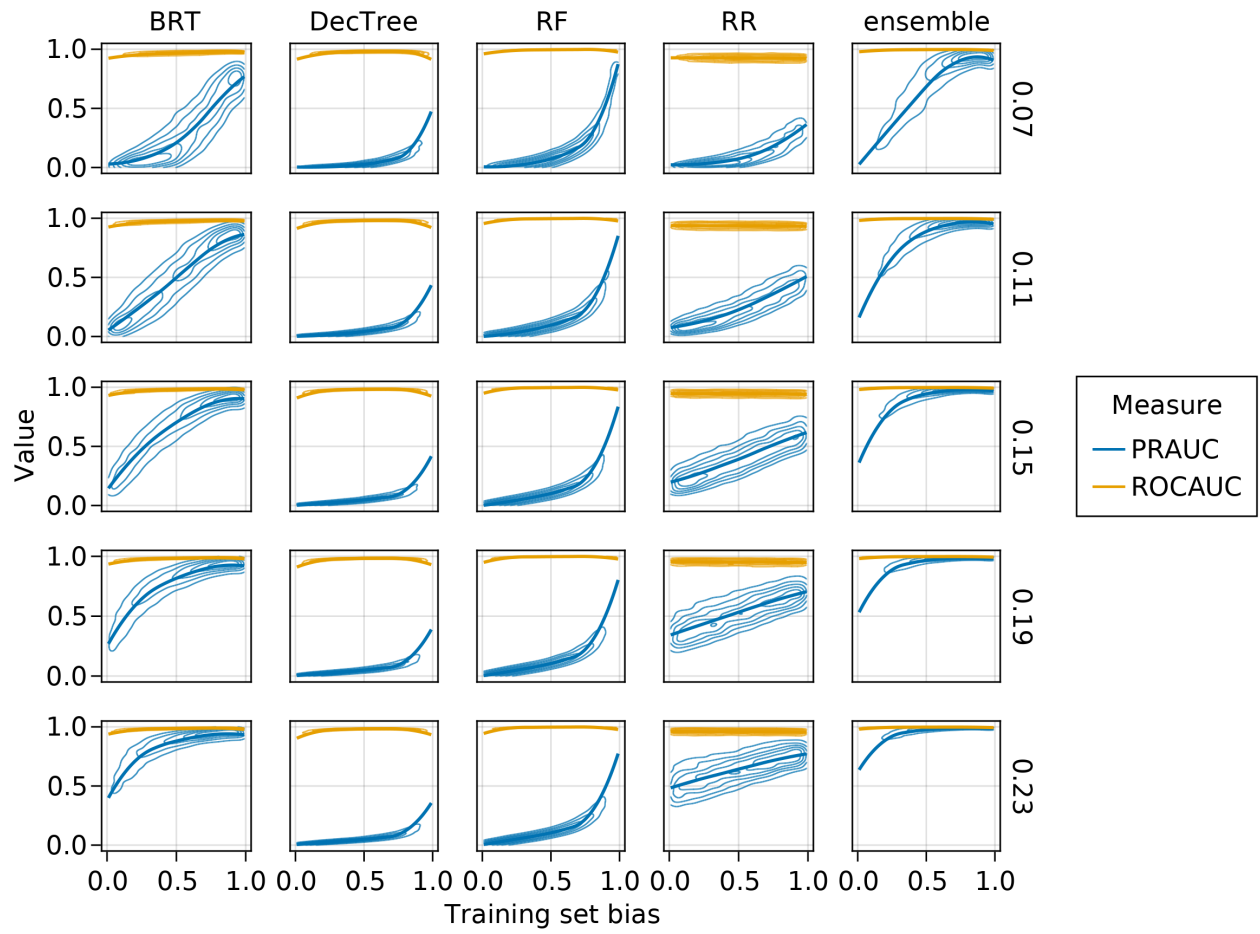
Figure 4: Response of ROC-AUC and PR-AUC to changes in the training set bias for a fixed connectance (rows). ROC-AUC is consistently high, and therefore not properly able to separate good from poor classifiers. On the other hand, PR-AUC responds to changes in the training set. As in fig. 3, training on more connected networks is easier.

Figure 5: Value of the optimal training set bias for the different models and measures evaluated here, over a range of connectances. Informedness was reliably maximized for balanced training sets, and kept this behavior across models. For other measures, larger connectances in the true network allowed lower biases in the training set. In a large number of cases, "over-correcting" by having training sets with more than half instances representing interactions would maximize the values of the model performance measures.
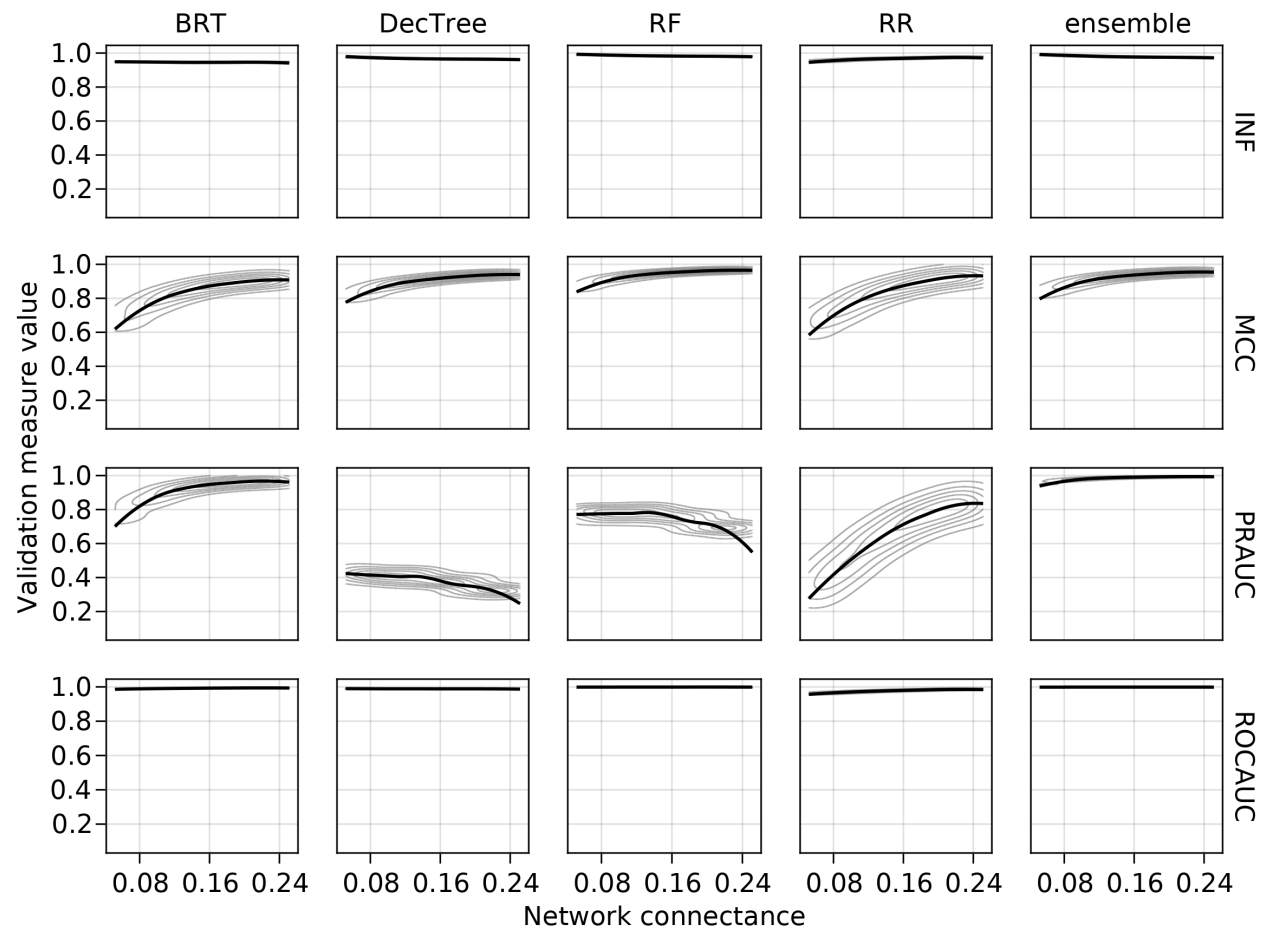
Figure 6: When trained on their optimally biased training set, most models were able to maximize their performance; this is not true for decision tree, which had a very low PR-AUC, and to some extent for ridge regression who had a slow increase with network connectance. The ensemble had a consistently high performance despite incorporating poor models.

Figure 7: Visualisation of the models predictions for one instance of a network prediction problem (shown in the "Dataset" panel). This figure reveals how inspecting the details of the prediction is important: indeed, although the performance measures hint at the fact that ridge regression is mediocre, this figure reveals that it is making predictions that correspond to a network with an entirely different topology (namely, nested as opposed to diagonal).